


Local Search and Optimization



Local Search Algorithms

- The search algorithms we have seen so far keep track of the **current state**, the “**fringe**” of the search space, and the **path** to the final state.
- In some problems, one doesn't care about a solution path but only the **orientation of the final goal state**
 - Example: 8-queen problem
- Local search algorithms operate on a **single state – current state** – and move to one of its **neighboring states**
 - Solution path needs not be maintained
 - Hence, the search is “local”

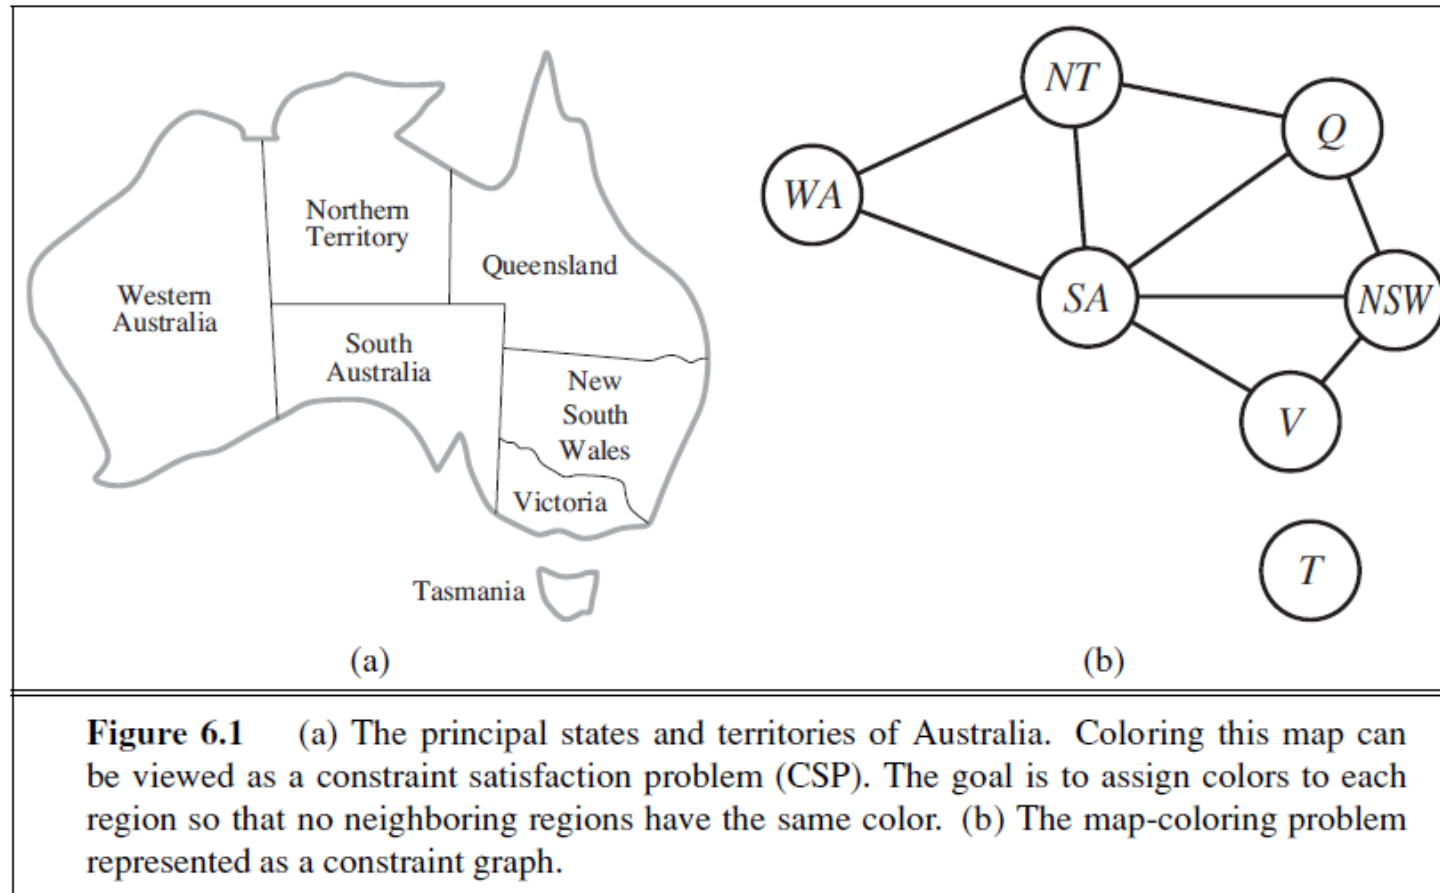
Example: Map-Coloring




- **Solutions** are **complete** and **consistent** assignments, e.g., WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green

Constraint graph

- It can be helpful to visualize a CSP as a **constraint graph**



Example




Local Search Algorithms

- When there is no need to memorize the path to the goal, just to find the goal.
- Cost of the path and finding optimal path to the goal is not required.
- Start from a single current state and check the neighboring states
- Consumes little memory
- Useful with large state space.

Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes



- Generally much faster and more memory efficient (but incomplete and suboptimal)

Local Search Methods – Applications

- Applicable when seeking Goal State & don't care how to get there. E.g.,
 - N-queens,
 - finding shortest/cheapest round trips
 - **(Travel Salesman Problem, Vehicle Routing Problem)**
 - finding models of propositional formulae **(SAT solvers)**
 - VLSI layout, planning, scheduling, time-tabling, . . .
 - map coloring,
 - resource allocation
 - protein structure prediction
 - genome sequence assembly

Local Search Algorithms

- In many optimization problems, the **path** to the goal is irrelevant; **the goal state itself is the solution**
- State space = set of configurations
 - Find a configuration satisfying your constraints, e.g., n-queens
 - Find the best possible state according to a given **objective function**
- In such cases, we can **use local search algorithms**
 - **Keeps a single "current" state**, and then shift states, but don't keep track of paths.
 - Use very limited memory
 - **Find reasonable solutions** in large state spaces.


Local Search Algorithm

- Local search algorithms work as follows:
 - Pick a “solution” from the search space and evaluate it. Define this as the current solution.
 - Apply a transformation to the current solution to generate and evaluate a new solution.
 - If the new solution is better than the current solution then exchange it with the current solution; otherwise discard the new solution.
 - Repeat steps 2 and 3 until no transformation in the given set improves the current solution

Local search


◆ Key idea (surprisingly simple):

1. Select (random) initial state (**generate an initial guess**)
2. Make **local modification** to **improve** current state
 - Evaluate current state and **move to other states**
3. Repeat Step 2 until goal state found (or out of time)



Local Search Algorithms for optimization Problems


- Local search algorithms are very useful for optimization problems
- systematic search doesn't work
- however, can start with a suboptimal solution and improve it
- Goal: find a state such that the objective function is optimized



**Minimize the number
of attacks**

Example: n-Queen

- Put N Queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts



Initial state ... Improve it ... using local transformations

Almost always solves n-queens problems instantaneously for very large n , e.g., $n =$
1 million

Example: n-Queen

- Put N Queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

Example: n-Queen

- Put N Queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

HEURISTIC:

h = # of pairs of queens that are attacking each other, either **directly** or **indirectly**

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18


Each number indicates h if we move a queen in its column to that square

12 (boxed) = best h among all neighbors;
select one randomly

Example: Travelling Salesman Problem

- **Travelling Salesperson Problem (TSP)**

- Suppose a salesman has five cities to visit and then must return home.
- The goal of the problem is to find the shortest path for salesman to travel, visiting each city, and then returning to the starting city.




Example: Travelling Salesman Problem




- A **Solution**: Exhaustive Search
 - **(Generate and Test) !!**
- The number of all tours is about $(n-1)!/2$
- If $n = 36$ the number is about:
56657398319307246483332566
87616000000000
- Not Viable Approach !!

Search for the travelling salesperson problem. Each arc is marked with the total weight of all paths from the start node (A) to its endpoint.




Example: Travelling Salesman Problem


Start with any complete tour, perform pairwise exchanges



Variants of this approach get within 1% of optimal very quickly with thousands of cities



Total Distance = 37km



Total Distance = 31km

Hill Climbing Algorithm

Gradient Ascent/Descent Algorithm



Hill Climbing

- Simple, general idea:
 - Start wherever
 - Repeat: move to the best neighboring state
 - If no neighbors better than current, quit
- What's bad about this approach?
 - Complete?
 - Optimal?



"Like climbing Everest in thick fog with amnesia"

- Hill climbing algorithm (a.k.a. greedy local search) uses a **loop** that continually moves in the direction of **increasing values** (that is uphill).
- It terminates when it reaches a peak where **no neighbor has a higher value**.

Hill-Climbing Algorithm

1. Pick a **random point** in the search space
2. Consider all the neighbors of the current state
3. Choose the neighbor with the best quality and move to that state
4. Repeat 2 to 4 until all the neighboring states are of lower quality
5. Return the current state as the solution state.

Hill Climbing Search

- Steepest Ascent Search, Greedy Local Search
- Iteratively maximize value of current state by replacing its successors state that has highest value , as long as possible.

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  
  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)  
  loop do  
    neighbor  $\leftarrow$  a highest-valued successor of current  
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE  
    current  $\leftarrow$  neighbor
```


Figure 4.2 The hill-climbing search algorithm, which is the most basic local search technique. At each step the current node is replaced by the best neighbor; in this version, that means the neighbor with the highest VALUE, but if a heuristic cost estimate h is used, we would find the neighbor with the lowest h .

Hill Climbing Search

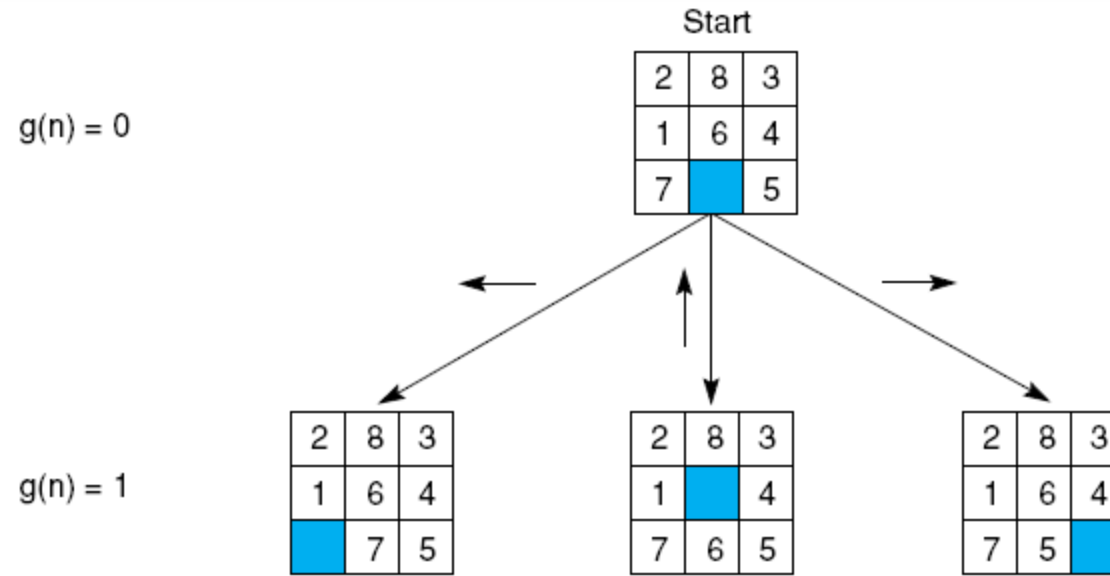
- Hill-climbing strategies expand the current state of the search and evaluate its children.
- The “best” child is selected for further expansion; neither its siblings nor its parent are retained.
- Because it keeps no history, the algorithm cannot recover from failures of its strategy.
- The algorithm does not maintain a search tree, so the data structure for the current node need only record the state and the value of the objective function.

- This is "hill climbing"
- We can use **heuristics** to guide "hill climbing" search.
- In this example, the **Manhattan Distance** heuristic helps us quickly find a solution to the 8-puzzle.

But "hill climbing has a problem..."



The heuristic f applied to states in the 8-puzzle.



Values of $f(n)$ for each state,

6

4

6

where:

$$f(n) = g(n) + h(n),$$


$g(n)$ = actual distance from n
to the start state, and


$h(n)$ = number of tiles out of place.

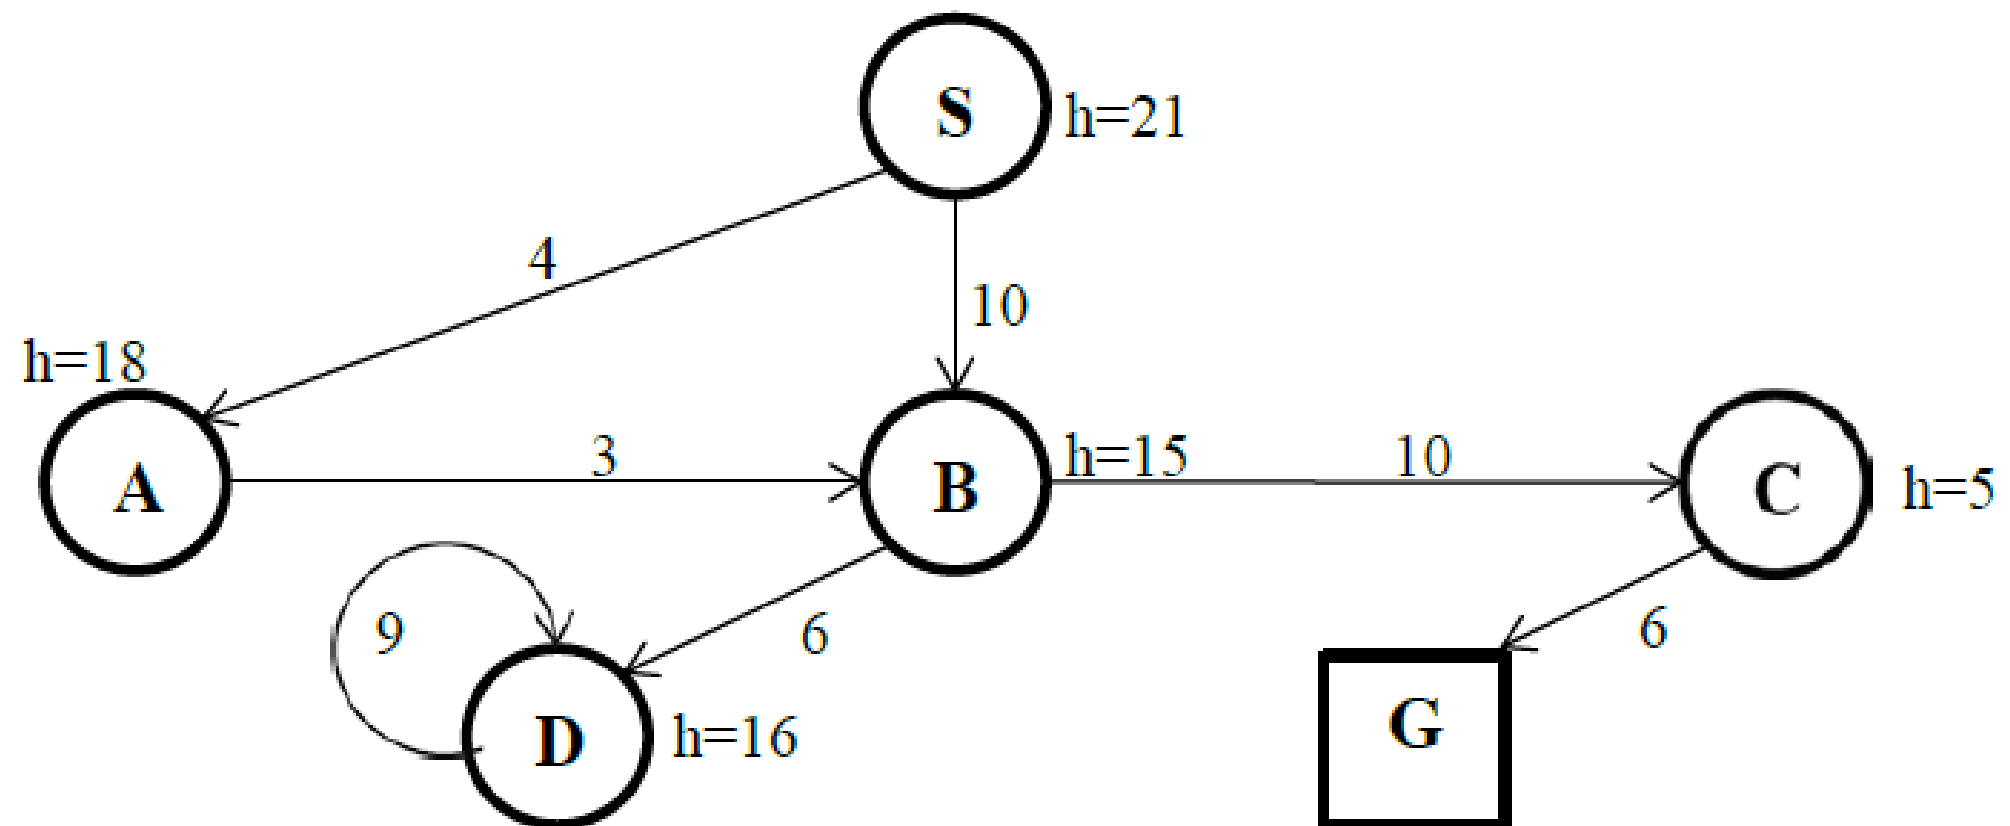
1	2	3
8		4
7	6	5


Goal

Hill climbing example (*minimizing* h)













Example of hill-climbing for the maze problem


- Consider the following maze:
- The problem is to get from the start node S to the goal node G, by moving horizontally and vertically and avoiding the black obstacles in the above maze.







- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.








- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.







- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.







- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.







- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.








- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.




Hill climbing





- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.



Hill climbing





- In the figures describing the steps of the algorithm
- each position will be denoted by a pair of numbers which are the horizontal and vertical coordinates (that is, for (x, y) , x denotes the horizontal coordinate and y denotes the vertical coordinate).
- The Manhattan distance between the points (x, y) and (z, t) is given by $d((x, y), (z, t)) = |x - z| + |y - t|$.



Example of hill-climbing for the maze problem

- The algorithm is stuck in (3,4) because all neighbours of (3,4) are worse than (3,4) (i.e., they have a higher Manhattan distance).




Example of hill-climbing for finding the minimum of a quadratic function


- Consider the function $f(x) = x^2 + 3x + 5$ defined on integer numbers in the interval $[-20, 20]$.
- Use the hill-climbing algorithm to find the function's minimum value on this interval.
- First define the search space, then the neighborhood. Try to find the minimum using starting point $x = 0$. Then try using starting point $x = -6$.

Example of hill-climbing for finding the minimum of a quadratic function

- The search space is the set $\{-20, -19, \dots, -1, 0, 1, 2, \dots, 19, 20\}$.
- The neighborhood for $x = 0$ is $\{-1, 0, 1\}$, for $x = -1$, $\{-2, -1, 0\}$ etc. In general for an x , the neighborhood is $\{x - 1, x, x + 1\}$.
 - $f(0) = 5$, $f(-1) = 3$, $f(1) = 9$, as we are looking for the minimum, hill climbing will select $x = -1$
 - $f(-1) = 3$, $f(-2) = 3$, $f(0) = 5$, so the algorithm will stop and give the answer $x = -1$, with the minimum value $f(-1) = 3$
- Starting from $x = -6$,
 - $f(-6) = 23$, $f(-7) = 33$, $f(-5) = 15$, so select $x = -5$.
 - $f(-5) = 15$, $f(-6) = 23$, $f(-4) = 9$, so $x = -4$ is selected.
 - $f(-4) = 9$, $f(-5) = 15$, $f(-3) = 5$, so $x = -3$ is selected.
 - $f(-3) = 5$, $f(-4) = 9$, $f(-2) = 3$, so $x = -2$ is selected.
 - $f(-2) = 3$, $f(-3) = 5$, $f(-1) = 3$, so $x = -2$ is returned as solution.

	Frontier (Hill Climbing(Expand	Explored
1	(S,9)	S	
2	(S-B,7)(S-C,10)(S-D,7)	B	
3	(S-B-E,1)	E	
4	(S-B-E-F,1) (S-B-E-G,0)	G	
5	(S-B-E-G,0)		





	Frontier (Hill Climbing(Expand	Explored
1	(S,12)	S	
2	(S-A,11)(S-D,10)	D	S
3	(S-D-S,12) (S-D-A,11) (S-D-E,7)	E	
4	(S-D-E-D,10) (S-D-E-B,6) (S-D-E-F,3)	F	
5	(S-D-E-F-C,3)		
6			

Problems with Hill Climbing in AI

Local Maximum

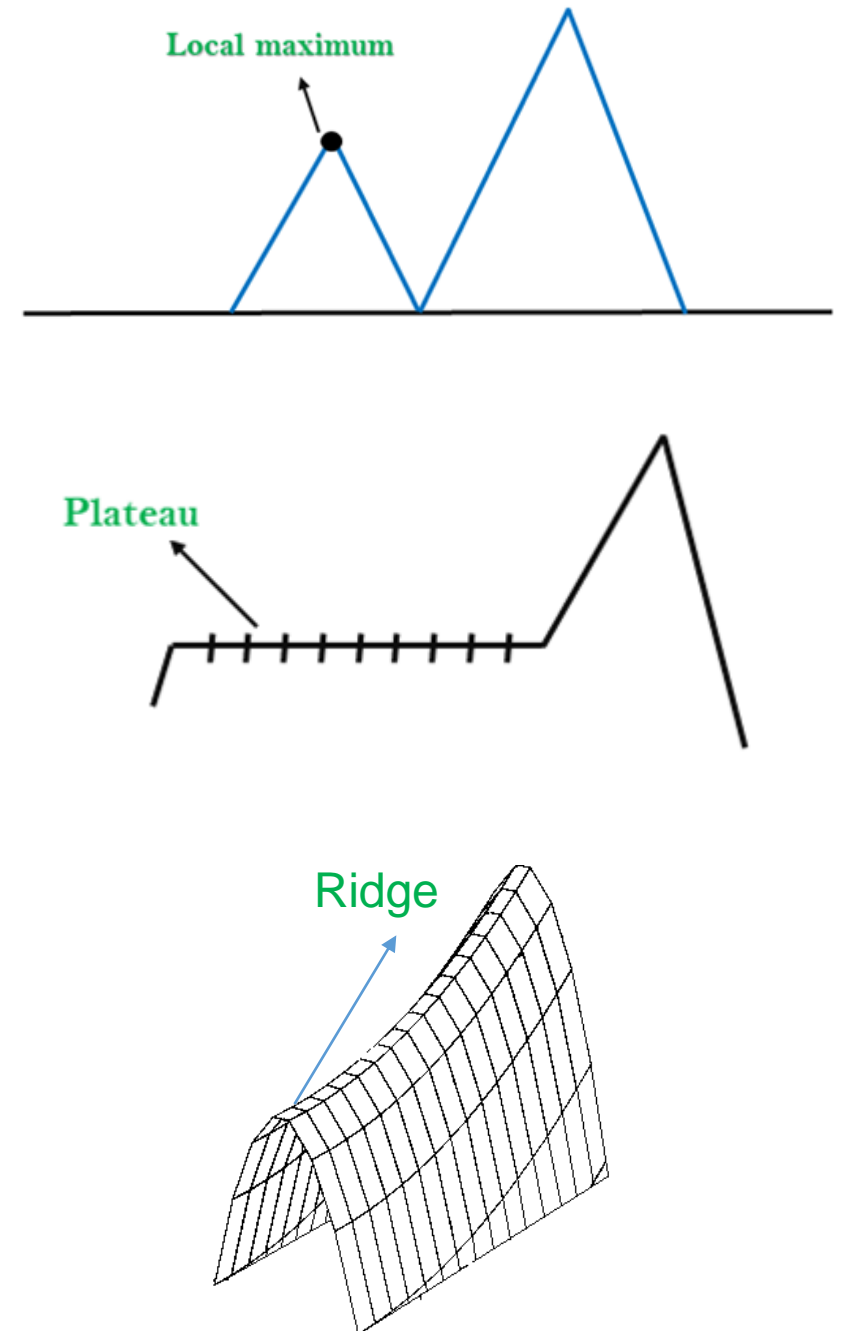
- A local maximum is a **peak state in the landscape** which is better than each of its neighboring states, but there is another state also present which is higher than the local maximum.

Plateau

- A plateau is the **flat area** of the search space in which all the **neighbor states of the current state contains the same value**, because of this algorithm does not find any best direction to move. A hill-climbing search might be **lost** in the plateau area.


Ridge

- A ridge is a **special form of the local maximum**. It has an area which is **higher than its surrounding areas**, but itself has a slope, and cannot be reached in a single move.




State Space Landscape

- A state space landscape is a **graph of states** associated with their costs
- A landscape has both “location” (defined by the state) and “elevation” (defined by the value of the heuristic cost function or objective function).



Hill Climbing



Hill-Climbing – Two Versions

- **Greedy Local Search**

- Systematically search the Neighbourhood.
- Accept the {first or best} cost decreasing move found.

- **Stochastic Local Search**


- Choose randomly from amongst the cost-decreasing moves found.

Hill-Climbing Problems

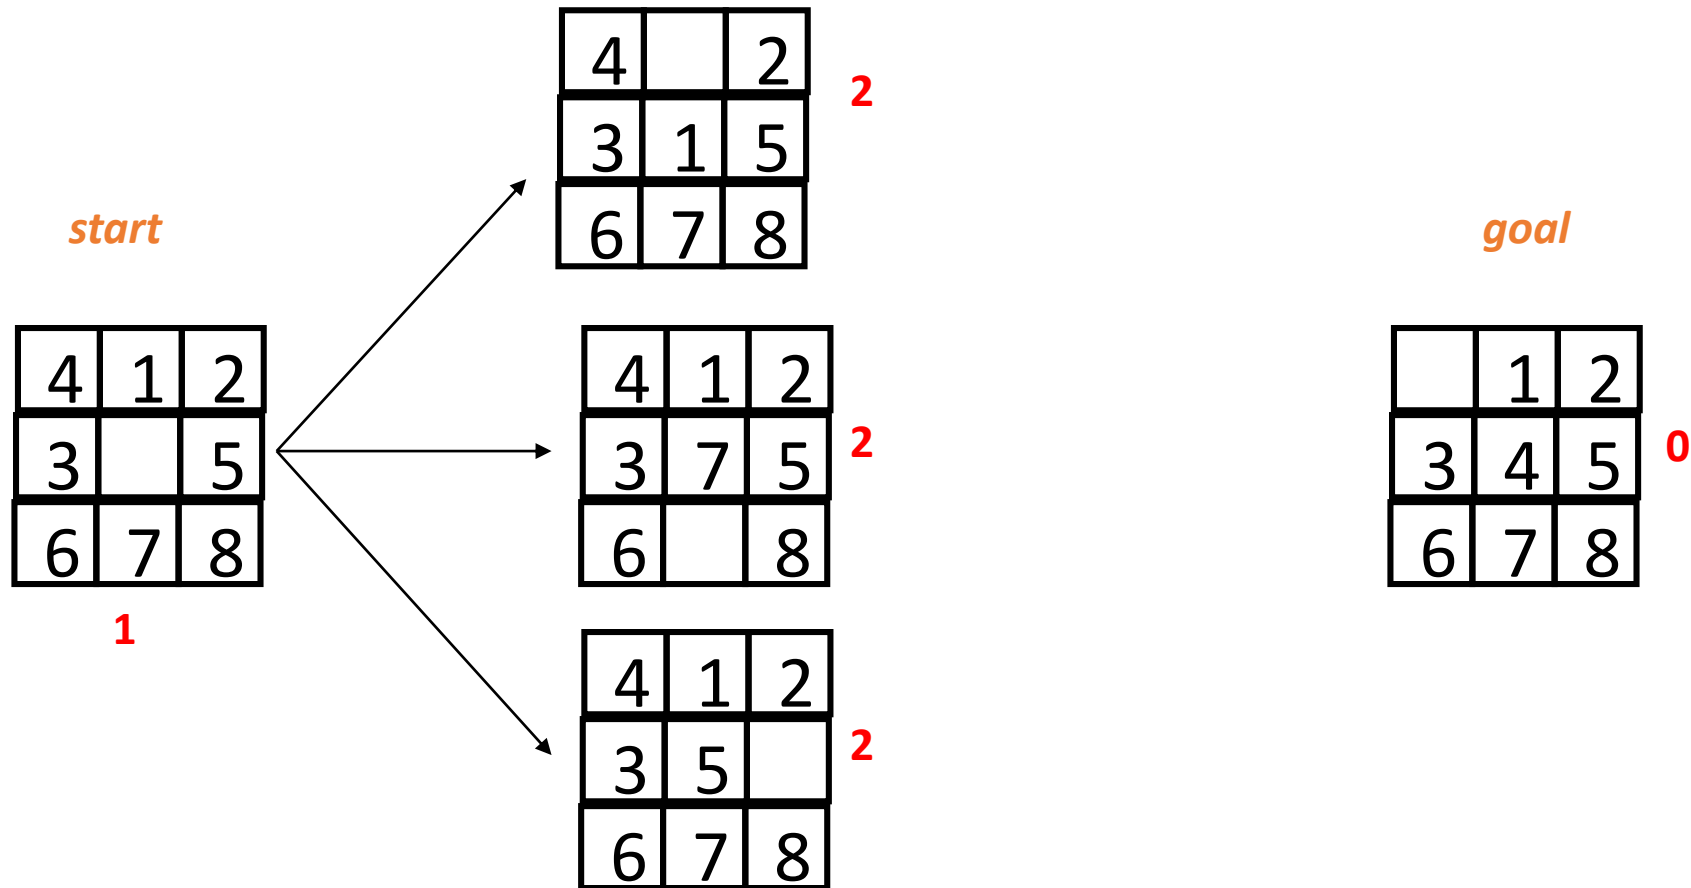
- Greedy Local Search: grabs a good neighbour state without thinking about where to go next
 - However, greedy algos do make good progress generally towards the solution

- Unfortunately, hill-climbing

- Can get stuck in **local maxima**
- Can be **stuck by ridges** (a series of local maxima that occur close together)
- Can be **stuck by plateau** (a flat area in the state space landscape)
 - Shoulder: if the flat area rises uphill later on
 - Flat local maximum: no uphill rise exists.





Example of a local "maximum"



Hill-Climbing Problems

- Local maxima/minima : local search can get stuck on a local maximum/minimum and not find the optimal solution



Local minimum

- Cure
 - + Random restart
 - + Good for Only few local maxima

Improvements

- **Random-restart Hill Climbing:** Selects a series of initial nodes randomly until the solution is found.
- Some problem spaces are great for hill climbing and others are terrible.

Variants of Hill Climbing

- Stochastic hill climbing:
 - chooses at random from among uphill moves
 - converges more slowly, but finds better solutions in some landscapes.
- First-choice hill climbing:
 - generate successors randomly until one is better than the current
 - good when a state has many successors
- Random-restart hill climbing:
 - conducts a series of hill climbing searches from randomly generated initial states, stops when a goal is found
 - It's complete with probability approaching 1

Variants

- In **simple hill climbing**, the first closer node is chosen, whereas in **steepest ascent hill climbing** all successors are compared and the closest to the solution is chosen. Steepest ascent hill climbing is similar to [best-first search](#), which tries all possible extensions of the current path instead of only one.
- [Stochastic hill climbing](#) does not examine all neighbors before deciding how to move. Rather, it selects a neighbor at random, and decides (based on the amount of improvement in that neighbor) whether to move to that neighbor or to examine another.
- **Random-restart hill climbing** iteratively does hill-climbing, each time with a random initial condition. The best is kept: if a new run of hill climbing produces a better than the stored state, it replaces the stored state.

Stochastic Hill Climbing

- Stochastic hill climbing chooses at random from among the uphill moves;
- The probability of selection can vary with the steepness of the uphill move.
- Stochastic hill climbing usually converges more slowly than steepest ascent, but in some state landscapes, it finds better solutions.
- Stochastic hill climbing is NOT complete, but it may be less likely to get stuck

First Choice Hill Climbing

- First-choice hill climbing implements stochastic hill climbing by generating successors randomly until one is generated that is better than the current state.
- This is a good strategy when a state has many of successors.
- First-choice hill climbing is also NOT complete,

Random Restart Hill Climbing

- Random-Restart Hill Climbing conducts a series of hill-climbing searches from randomly generated initial states, until a goal is found.
- Random-Restart Hill Climbing is complete if infinite (or sufficiently many tries) are allowed.
- If each hill-climbing search has a probability p of success, then the expected number of restarts required is $1/p$.
- The success of hill climbing depends very much on the shape of the state-space landscape:
- If there are few local maxima and plateau, random-restart hill climbing will find a good solution very quickly.
- On the other hand, many real problems have many local maxima to get stuck on.

Nature-Inspired Heuristic Algorithms

- **Genetic Algorithm** (evolution)
- Neural Network (artificial neural network)
- Simulated Annealing (metal annealing)
- Tabu Search (animal's brain)
- Ant Colony Optimization
- Bee Colony Optimization